

DEAD LOCKS



Resource Allocation

Sharable Resources

Resources that can be shared among number of processes

- Memory, Hard Disk etc..

Non Sharable resources

Resources that cannot be shared among number of processes

- Printer, Tape Drive, CDROM Drive etc...

A resource type can have more than one instance

Resource Allocation

Under normal operation, a resource allocations proceed like this::

1. Request a resource (suspend until available if necessary).
2. Use the resource.
3. Release the resource.

Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

➤ Mutual Exclusion

- Only one process at a time can use a resource.

➤ Hold and Wait

- A process holding at least one resource is waiting to acquire additional resources held by other processes.

Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

➤ No Resource Preemption

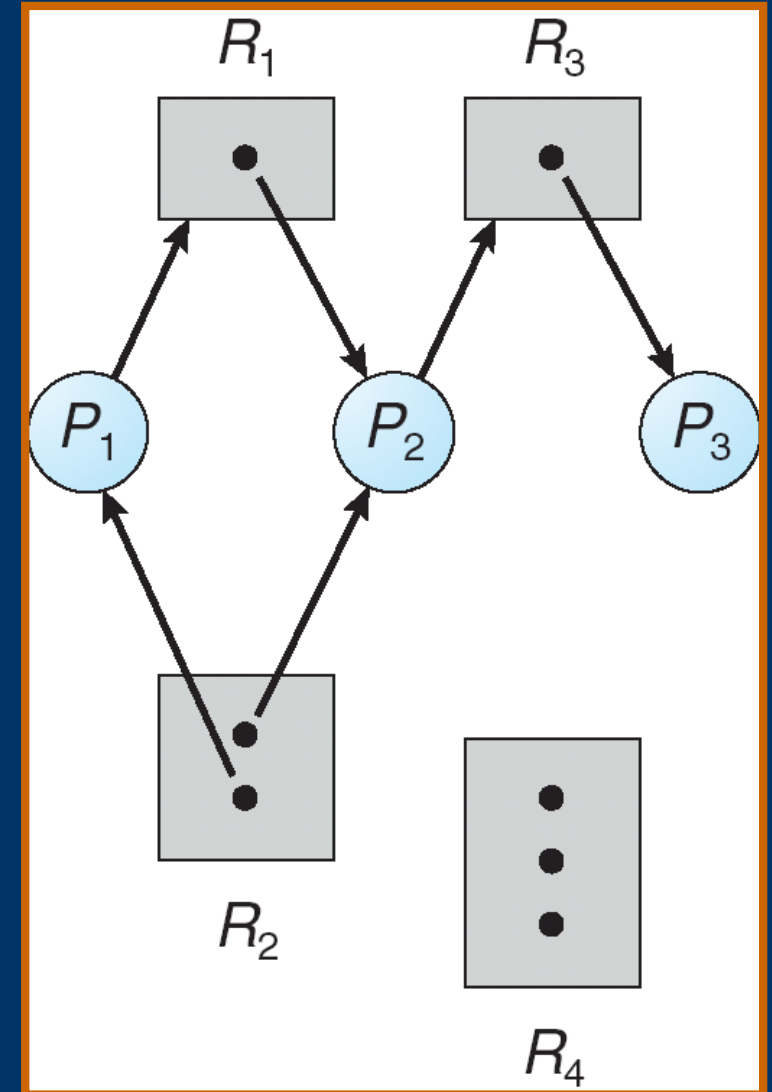
- A resource can be released only voluntarily by the process holding it after that process has completed its task.

➤ Circular Wait

- There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Directed Resource Allocation Graph

- **DRAG** is a Graphical representation showing the relationship between processes and resources.
- Helps in detecting deadlocks.




Resource-Allocation Graph


A set of vertices V and a set of edges E .

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- Request edge - directed edge $P_i \rightarrow R_j$
- Assignment edge - directed edge $R_j \rightarrow P_i$

Resource Allocation Graph




Resource Allocation Graph

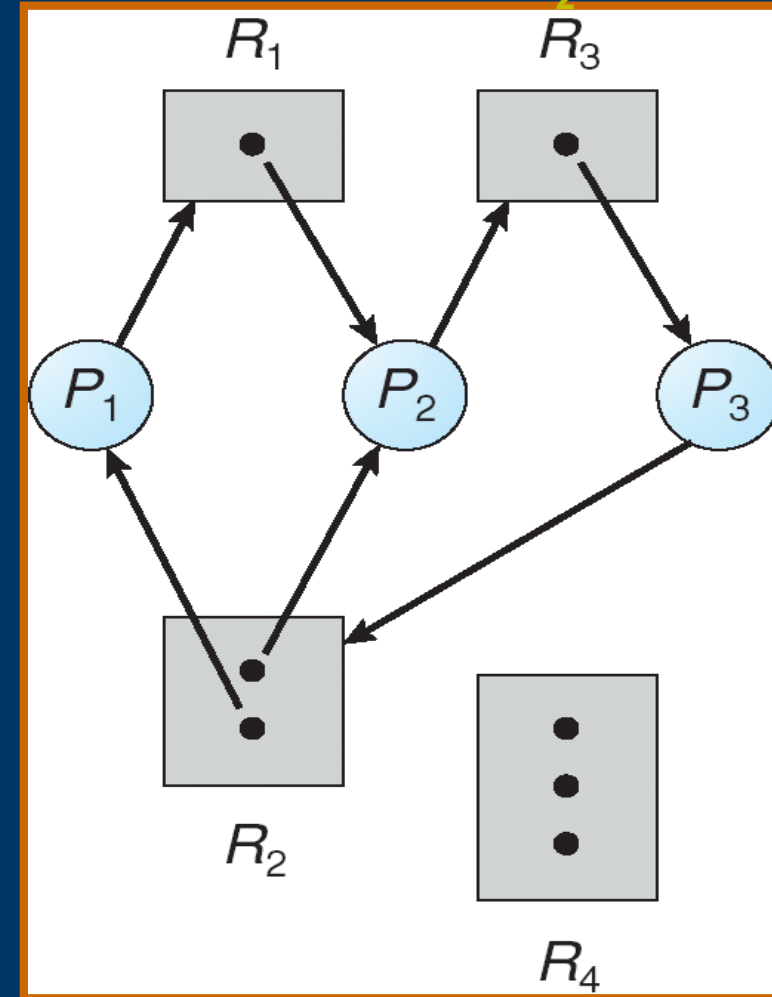


Resource Allocation Graph With A Deadlock


Before P_3 requested an instance of R_2



After P_3 requested an instance of R_2



Graph With A Cycle But No Deadlock



Process P_4 may release its instance of resource type R_2 . That resource can then be allocated to P_3 , thereby breaking the cycle.

Relationship of cycles to deadlocks

- If a resource allocation graph contains no cycles
⇒ No deadlock
- If a DRAG contains a cycle and if only one instance exists per resource type
⇒ Deadlock
- If a DRAG contains a cycle and if several instances exists per resource type
⇒ Possibility of deadlock

Deadlock Strategies



- Ignoring Deadlock

- Detecting Deadlock

- Recovering From Deadlock State

- Deadlock Prevention

- Deadlock Avoidance

Deadlock Avoidance

- **Deadlock prevention**
 - Low device utilization
 - Low system throughput
- **Deadlock Avoidance**
 - Additional information about how resources are to be requested.

Deadlock Avoidance

- Each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- **Resource-allocation state** is defined by the number of available and allocated resources, and the maximum demands of the processes

Deadlock Avoidance

- Allows more concurrency than prevention
- Deadlock Avoidance Algorithm
 - Do not **start** a process if its total demand might lead to deadlock.
 - Do not grant an **incremental** resource request if this allocation could lead to deadlock.

Deadlock Avoidance

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.
- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$


Deadlock Avoidance

- If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished.
- When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
- When P_i terminates, P_{i+1} can obtain its needed resources, and so on.

Basic Facts

- If a system is in safe state \Rightarrow no deadlocks
- If a system is in unsafe state \Rightarrow possibility of deadlock.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

Safe, Unsafe, Deadlock State



Banker's Algorithm

- Tentatively grant each resource request.
- Analyze resulting system state to see if it is “safe”.
- If safe, grant the request.
- If unsafe refuse the request (undo the tentative grant)
- Block the requesting process until it is safe to grant it.


Data Structures for the Banker's Algorithm

- Let n = number of processes, and m = number of resources types.
- **Available** : Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- **MAX** : $n \times m$ matrix. If $Max [i,j] = k$, then process P_i may request at most k instances of resource type R_j

Data Structures for the Banker's Algorithm

- **Allocation:** $n \times m$ matrix. If $\text{Allocation}[i, j] = k$ then P_i is currently allocated k instances of R_j
- **Need:** $n \times m$ matrix. If $\text{Need}[i, j] = k$, then P_i may need k more instances of R_j to complete its task
$$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$$

Available [m]



- Vector of length m .
- If $\text{available}[j] = k$, there are k instances of resource type R_j available.

MAX [n][m]

	Printer	Tape	Plotter
P0	7	5	3
P1	3	2	2
P2	9	0	2
P3	2	2	2
P4	4	3	3

- If $Max [i, j] = k$, then process P_i may request at most k instances of resource type R_j

Allocation [n][m]

	Printer	Tape	Plotter
P0	0	1	0
P1	2	0	0
P2	3	0	2
P3	2	1	1
P4	0	0	2

- If Allocation[i, j] = k then P_i is currently allocated k instances of R_j

	Printer	Tape	Plotter
P0	7	5	3
P1	3	2	2
P2	9	0	2
P3	2	2	2
P4	4	3	3

	Printer	Tape	Plotter
P0	0	1	0
P1	2	0	0
P2	3	0	2
P3	2	1	1
P4	0	0	2

$$Need [i, j] = Max [i, j] - Allocation [i, j]$$

Need [n][m]

	Printer	Tape	Plotter
P0	7	4	3
P1	1	2	2
P2	6	0	0
P3	0	1	1
P4	4	3	1

- If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively

Initialize:

(a) *Work* = *Available*

(b) For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then

$Finish[i] = \text{false}$; otherwise, $Finish[i] = \text{true}$

2. Find an index i such that both:

(a) $Finish[i] == \text{false} \ \&\& \ Need_i \leq Work$

If no such i exists, go to step 4

Safety Algorithm

3. $Work = Work + Allocation_i$

$Finish[i] = true$

go to step 2

4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in deadlock state.

Safety Algorithm

Work = Available, Finish[i]='F'



5 - Processes

10 - Printers, 5 - Tapes, 7 - Plotter